

MICKAEL ZEITOUN

AI Engineer | Data Science M.Sc.

CONTACT

-  mikizit97@gmail.com
-  058-7310-897
-  github.com/mikigit97
-  [linkedin/mickael-zeitoun](https://linkedin.com/mickael-zeitoun)

TECHNICAL SKILLS

Machine Learning

Python, PyTorch, Hugging Face
Transformers, LLMs & LoRA/PEFT
fine-tuning, foundation models, NLP
(TF-IDF, DistilBERT), scikit-learn,
feature engineering

Software & Data

Docker, Linux, SQL, REST/HTTP APIs,
Slurm/cluster GPU workflows, data
preprocessing, experiment design,
model evaluation

Programming

Python, SQL, Java, R

Tools

Git, VS Code, Ollama, TensorBoard,
Jupyter, Colab

LANGUAGES

English: Fluent

Hebrew: Fluent

French: Native

EDUCATION

M.Sc. Data Science

Ben-Gurion University
Meitar Excellence Program
Graduated

B.Sc. Data Engineering

Ben-Gurion University
Completed
Data Structures (97), Computational Models
(99)

PROFESSIONAL SUMMARY

AI Engineer with a Data Science M.Sc. focused on applied LLM and NLP work: parameter-efficient fine-tuning (LoRA/PEFT), local LLM-agent pipelines, and classical-to-transformer NLP benchmarking. Comfortable running GPU training jobs on a compute cluster and combining language models with other techniques, such as graph analysis, to solve applied problems. Backed by a Data Engineering B.Sc. for production-minded software skills.

RESEARCH EXPERIENCE

LLM Fine-Tuning with LoRA (Sherlock Holmes corpus)

Used LoRA, a parameter-efficient fine-tuning method, to adapt a small language model (Qwen2.5-0.5B) to a literary corpus: continued pre-training cut perplexity 18.05 to 16.54, and supervised fine-tuning raised QA F1 from 0.011 to 0.291. Ran 10+ ablations as GPU jobs on a SLURM cluster.

LLM-Agent User-Behavior Simulation

Built a five-phase simulation framework using LLM agents (Ollama) and MovieLens-32M to generate labeled interaction data for three user types, feeding a downstream recommendation-system classifier.

Large-Scale Network & Graph Analysis

Analyzed real networks with centrality metrics, PageRank, and node2vec embeddings for directed link prediction. Combined centrality features with a locally hosted LLM (Ollama) over HTTP to detect managers in the Enron email network.

Tweet Device Classification (NLP)

Engineered temporal, stylometric, and TF-IDF features and benchmarked five model families, including a fine-tuned DistilBERT, under 5-fold stratified cross-validation with grid search. Best F1: 0.862.

EXPERIENCE

CyberBGU - AI Integration Research

2024-now
Partnered with RAD to integrate AI-driven optimization into Industrial Control Systems (ICS), targeting network load efficiency and operational performance.

Samsung - Technical Support

2020
Diagnosed and resolved technical issues under time pressure and communicated solutions clearly to customers.

MILITARY SERVICE

Israeli Air Force - F-16 Aeronautics Technician (Sergeant Major)

Diagnosed mechanical and electronic faults in real-time, safety-critical systems.

VOLUNTEERING

PERACH mentoring | Special-needs youth support | Holocaust survivor visits